

Smart Money on Polymarket: A Behavioral Anatomy of 273 Top Prediction-Market Traders

Empirical Study

Wallet-level data: `polydata.pro`, snapshot of 21 April 2026

April 2026

Abstract

We assemble a wallet-level panel of the 273 most-analyzed traders on [Polymarket](#) and ask a single question: *how do prediction-market “smart money” wallets actually make money?* Using publicly available analytics aggregated by `polydata.pro`, we observe each trader’s full trade history, risk-adjusted performance, hour-of-day activity, DCA intensity, market specialization and a thirteen-dimension bot-detection fingerprint. Three first-order findings emerge. **(i) Profits are extraordinarily concentrated:** the top decile of profitable wallets captures 52% of aggregate positive PnL (\$98M of \$189M); the top quintile captures 89%. **(ii) The dominant strategy is not arbitrage but *resolution-edge holding*:** median *event*-level win rate is 93% while median *trade*-level win rate is only 58%, consistent with heavily-staggered DCA entries (median 15 entries per market) into binary outcomes that the holder believes are mispriced relative to ultimate resolution. **(iii) 67% of analyzed wallets are at least partially automated**, but pure bots earn the lowest median PnL (\$33K) of any group; the most profitable cohort are “likely human” fast-interface traders (\$84K median). Politics is the dominant profit pool—98% of politics specialists are profitable, with median PnL \$172K—while sports attracts the largest population (45%) but a less skewed payout distribution. K-Means clustering on behavioural features yields six distinguishable strategy archetypes, with “Heavy-DCA Holders” producing the highest median profit and “HFT-like Bots” the lowest. We argue that “smart money” on Polymarket is best understood not as alpha-extracting arbitrage but as *patient capital with strong directional priors that exploits a slow-resolving information environment*.

1 Introduction

[Polymarket](#) is the largest decentralised prediction market by traded volume. Built on Polygon [[Polymarket, 2025](#)] and collateralised in USDC, it allows participants to buy and sell shares of binary outcomes priced between 0¢ and 100¢; correct shares pay out at 100¢ on resolution. Prior literature on prediction markets emphasises information aggregation [[Wolfers and Zitzewitz, 2004](#), [Arrow et al., 2008](#), [Manski, 2006](#), [Atanasov et al., 2017](#)] and behavioural biases such as the favourite–longshot effect [[Griffith, 1949](#), [Snowberg and Wolfers, 2010](#), [Thaler and Ziemba, 1988](#)]. Comparatively little work has examined the *individual traders* who systematically extract profits from these markets, despite the on-chain settlement of Polymarket making such analysis uniquely tractable.

This paper takes the smallest-possible-step empirical approach. We download the full per-wallet analytics published by the third-party service `polydata.pro` [[Polydata, 2026](#)] for every wallet they have “analyzed” (a sample of 277 wallets selected by Polydata’s pipeline based on activity and notability; 273 deliver complete pages and form our analysis sample). Each wallet’s page exposes more than sixty behavioural and performance metrics derived from on-chain trade

history, including a 0–100 “Bot Score” built on thirteen detection dimensions. We treat the resulting cross-section as a complete census of “smart money” suspects on the platform and ask:

1. How concentrated are profits, and what is the empirical shape of the smart-money payoff distribution?
2. What *strategy* do successful wallets follow? Is it arbitrage across correlated outcomes, statistical edge from polling and macro data, counter-party-providing market making, or something else?
3. How prevalent is automation, and do bots or humans extract more profit per dollar of risk?
4. Which markets (Politics, Sports, Crypto, Weather, etc.) reward participation, and which entrench existing skill?

The headline answer to (2) is provided in the abstract: a 93% event win rate combined with only 58% trade win rate is hard to reconcile with arbitrage or with any short-horizon scalping story; it is precisely what one would expect from *DCA accumulation into a probabilistic view, held to resolution*—more akin to value investing with a short half-life than to high-frequency trading.

Roadmap. Section 2 describes data collection and the panel. Section 3 documents the unconditional distribution of PnL, volume, win rates and tier composition. Section 4 analyses the bot-population in depth: the marginal contribution of bot behaviour, the trade-gap fingerprint, and the share of volume vs PnL that automation explains. Section 5 runs K-Means on behavioural features and labels six strategy archetypes. Section 6 contrasts profitability across market verticals. Section 7 drills down on DCA, concentration, holding style, and risk-adjusted performance. Section 8 synthesises the findings into an answer to the central question—*what is the source of smart-money edge?*— and Section 9 concludes.

2 Data and Methodology

2.1 Source and scraping

The polydata.pro public site exposes server-side-rendered Next.js pages whose initial state contains a structured JSON payload (`initialData`) with all metrics displayed in the UI. We scraped the leaderboard index (<https://polydata.pro/traders>) to obtain the universe of 277 “analyzed” wallets, then fetched each wallet detail page (<https://polydata.pro/traders/<address>>) at a polite concurrency of 6 req/s with retry. Three wallets returned an `$undefined` sentinel even after multiple retries (server-side analysis evidently in progress) and are dropped, leaving $N = 273$ traders in the analytic sample. Raw HTML, parsed JSON and a flat CSV are released alongside the paper.

2.2 Panel structure

For each wallet we extract:

- *Identity*: nickname, wallet address, X (Twitter) handle, account creation date, verified status.
- *Performance*: net PnL (USD), realised PnL, total trading volume, daily volume, portfolio value at snapshot, leaderboard rank, per-category and per-period rankings.
- *PnL composition*: number of closed positions, wins/losses, trade-level and event-level win rates, profit factor, average win/loss, top-10 winning and top-10 losing positions with titles.

- *Risk*: Sharpe and Sortino on daily PnL, Herfindahl index of position concentration (HHI), top-1 and top-5 market shares, Kelly % and EV per dollar.
- *Behaviour*: trades per active day, markets per active day, 24-hour and day-of-week activity histograms, peak-hour UTC, median and mean trade gap (seconds), % of trades at night, position size coefficient of variation, % of trades sold vs held to resolution.
- *Bot fingerprint*: a Polydata-supplied 0–100 Bot Score plus ≤ 13 qualitative signals (e.g., “51% of trades within 5 s”, “24/24 active hours”, “size CV 3.5”), and a four-class verdict (Human, Likely human, Semi-bot, Bot).
- *Style tags*: pseudo-binary badges (Sniper, Whale, Diamond Hands, Lottery Player, Contrarian, etc.), DCA intensity, YES bias %, low/mid/high probability-band exposure, weather/sports specialization %.
- *Time-series*: full daily equity curve (from first trade to the snapshot of 21 April 2026); per-session start, end, duration and trade count for ≤ 50 recent sessions.

The flattened wide table contains 273×106 cells, supplemented with auxiliary tables of equity curves (18 848 rows), top winning and losing positions ($2\,542 + 2\,336$ rows), hour distributions (6 552 rows), session records (13 004 rows), and bot signals (1 775 rows).

2.3 Selection-bias caveat

The Polydata sample is not random. Wallets are added to the analytics pipeline based on activity, notability, and likely user request. The sample therefore over-represents profitable, high-volume, and high-profile wallets. Concretely, 83% of our panel is profitable, versus a much lower share of profitable participants estimated for Polymarket as a whole. All distributional claims below are conditional on this selection. Where we discuss bot prevalence, win-rate asymmetries or strategy archetypes, the patterns are still informative about the upper tail of trading skill on the platform—which is arguably the population of greatest interest.

3 Descriptive Statistics

3.1 The shape of profitability

Across $N = 273$ analyzed wallets, aggregate net PnL is **\$177.3 M**, against **\$9.5 B** of cumulative trading volume—an aggregate 1.87% return on volume. Mean PnL per trader is \$649K but median is only \$69K, foreshadowing the extreme right skew that defines the sample. Figure 1 plots the distribution of $|\text{PnL}|$ on a log axis with positive and negative wings overlaid. Positive PnL spans six orders of magnitude, from \$1 up to Theo4’s \$22.05M; negative PnL is much more compressed—only 46 wallets (16.8%) lose money and the worst loss is \$4.5M.




Figure 1: Distribution of net PnL across 273 analyzed Polymarket traders. Positive (green) and absolute negative (red) PnL plotted on a common log axis. The right tail extends six orders of magnitude.

Volume distribution (Figure 2) is even more extreme: a handful of wallets clear several hundred million USD in volume each. Combined, these two distributions yield the Lorenz-style profit concentration in Figure 3: **the top 20% of profitable wallets capture 89% of aggregate positive PnL**, and the top decile alone captures 52%. By any reasonable threshold, “smart money” on Polymarket is a tiny minority.




Figure 2: Distribution of total trading volume.

Figure 3: Lorenz-style cumulative profit concentration.

The top twenty wallets are listed in Figure 4, coloured by bot classification. The single largest profit (\$22M, Theo4) and the #2 wallet (Fred9999, \$16.6M) are both classified Semi-bot and specialise in Politics. The first pure “Bot” classification appears at #3 (kch123, \$11.9M, Sports) and the first pure “Human” enters the top twenty only at rank fourteen.




Figure 4: Top twenty wallets by realised net PnL, coloured by Polydata’s bot classification.

3.2 Smart-Score tiers

Polydata produces a 0–100 “Smart Score” aggregating eight performance components (Win Rate, Profit Factor, Sharpe, Sortino, EV, HHI, Profit Days, DD/Peak). Figure 5(a) shows that 36% of analyzed traders attain “Elite” status (≥ 90) while 19% fall into the bottom two tiers despite being analyzed. Median PnL by tier (panel b) is monotone: Elite traders earn an order of magnitude more than Average traders.

Figure 5: Smart-Score tier distribution and profitability




Figure 5: Smart-Score tier distribution and median PnL by tier.

3.3 Win rates: trade-level vs event-level

The single most diagnostic statistic in our panel is the gap between *trade-level* and *event-level* win rates. The median trader wins only **58%** of individual trades but resolves **93%** of *events* (multi-trade positions on a single underlying market) profitably. Figure 6 plots the two against

each other, coloured by signed log PnL. The mass of points sits well above the diagonal: most wallets accept many small losing trades within an event, but ultimately exit (or hold to resolution) on the winning side. This is the defining footprint of *accumulating-into-a-view-and-holding* as opposed to scalping.

Figure 6: Event vs trade win rates (event $\gg$ trade $\Rightarrow$ DCA winners)




Figure 6: Event-level versus trade-level win rates. The dashed line is the diagonal. Mass above the diagonal indicates DCA-style accumulation.

3.4 Risk-adjusted performance and tenure

Sharpe ratios on daily PnL (Figure 7) are mostly modest (median 0.46); only 13.6% of analyzed traders achieve a Sharpe above 1. Conditional on bot class, distributions are broadly similar in location but humans show fatter tails on both sides. Account-age versus PnL (Figure 8) reveals essentially no monotone relationship within the analyzed sample: short-tenure wallets (< 60 days) include several of the largest profit producers, suggesting that capital and information matter more than experience curves.




Figure 7: Sharpe distribution by bot class.




Figure 8: Account age vs PnL (colour = bot score).

4 Bot Detection and Its Profit Implications

4.1 Population breakdown

Polydata classifies 64 wallets (23.4%) as outright Bot, 120 (43.9%) as Semi-bot, 68 (24.9%) as Likely human, and 21 (7.7%) as Human—so **67.4% of analyzed wallets are at least partially automated**. The continuous Bot Score (Figure 9) is approximately bimodal, with mass at 0 (clear humans) and ~ 60 –80 (clear bots), and a wide hump in the middle (semi-bots).




Figure 9: Distribution of Polydata’s Bot Score across analyzed traders.

4.2 Do bots make more money?

A naive prior is yes—automation should win on speed. The data disagrees. Figure 10 shows PnL distributions by bot class on a symmetric log scale. Median PnL by class is \$33K (Bot), \$74K (Semi-bot), \$84K (Likely human), \$38K (Human). Pure bots earn the *lowest* median; the most profitable cohort is the “Likely human” fast-interface group, which uses scripts and tooling but retains a human in the loop. The interpretation is straightforward: bots provide *execution quality*, not *view quality*. Without an edge, fast hands lose the bid–ask to the same degree as slow ones.




Figure 10: PnL distribution by bot classification (symmetric log).

4.3 Trade-gap fingerprint

Figure 11 plots the histogram of median time between trades, by bot class. Pure bots cluster below 5 seconds (with a spike near 1 second indicating market-data-driven loops); semi-bots straddle 10–100 s; humans show median gaps of minutes to hours. A simple rule—“median trade gap below 5 s”—flags 24 wallets (8.8%) and is among the strongest single bot indicators. The companion frequency vs coverage plot (Figure 12) shows that bots also occupy a distinctive top-right region: high trade volume *and* broad market coverage.




Figure 11: Median trade-gap distribution.




Figure 12: Trade frequency vs market coverage.

4.4 Hour-of-day fingerprint

Bots are temporally undifferentiated. Figure 13 shows the 24-hour trade-volume heatmap by bot class. Humans concentrate trades in the US business afternoon (16–22 UTC) and trail off overnight; likely-human traders extend the active window; semi-bots and bots operate uniformly across the clock. 220 of 273 wallets (80.6%) are flagged for “24/24 active hours” and this is the single most common bot signal in the dataset (Figure 15).




Figure 13: 24-hour trade volume share by bot class.

4.5 The volume-profit divergence

Figure 14 decomposes the bot-class share of (a) trader count, (b) total trading volume, and (c) aggregate PnL. Bots represent 23% of traders but 25% of volume and only 15% of profits; semi-bots are 44% of traders, 57% of volume and 62% of profits; humans collectively are 33% of traders but 18% of volume and 24% of profits. **Bots churn at a loss-of-edge rate; humans extract more profit per dollar of volume.**




Figure 14: Bot class share of trader count, volume, and aggregate PnL.




Figure 15: Most common bot-detection signals across analyzed traders.

5 Strategy Clustering

We cluster the 273 wallets on a 17-dimensional behavioural feature vector that includes trades-per-day, markets-per-day, median trade gap, night-trading %, active hours, average DCA entries, YES bias, exposure to low/mid/high probability bands, sell rate, event and trade win rates, HHI, top-5 share, average session minutes and trades per session, and weather specialization. Highly skewed features are $\log(1+x)$ -transformed, then all features are standardised. K-Means with $K = 6$ produces an interpretable taxonomy:

1. **Heavy-DCA Holder** — median ~ 100 entries per market, near-zero sell rate, holds to resolution. Highest median PnL.
2. **Mixed-Event Sniper** — moderate DCA, broad market coverage, fast trade gaps. The largest single cluster.
3. **HFT-like Bot** — ≥ 500 trades/day, sub-30 s median gap, near-uniform 24-hour activity. Lowest median PnL.
4. **Active Trader (Sells)** — material sell rate, lower DCA, shorter holding periods.
5. **Long-tail Lottery** — exposure overweighted on low probabilities ($< 20\%$), winner-take-all payoff style.
6. **Weather Specialist** — $\geq 40\%$ of activity in weather markets, otherwise heterogeneous.

Figure 16 shows the PCA projection coloured by cluster; Figure 17 renders each cluster as a normalised radar of its eight most distinguishing dimensions. Figure 18 plots PnL distributions per cluster: *Heavy-DCA Holders* dominate, while *HFT-like Bots* cluster around break-even. Figure 19 crosses each cluster against its bot composition, confirming that “HFT-like Bot” is unanimously automated whereas the high-PnL “Heavy-DCA Holder” cluster is dominated by Likely-human and Semi-bot classes.




Figure 16: PCA projection of behavioural features.




Figure 17: Behavioural fingerprint per cluster.




Figure 18: PnL by cluster.




Figure 19: Cluster vs bot-class composition.

6 Category Specialization

6.1 Where the money is

Sports markets attract the most participants (45%) but Politics is where smart money concentrates and where the outsized profits are booked. Aggregate PnL by primary category is dominated by Politics (\$108.9M of the \$177.3M total), followed by Sports (\$55.6M) and Crypto (\$9.9M). Politics specialists are profitable in 98.4% of cases (median PnL \$172K, mean \$1.73M), versus 76.6% for Sports and 79.2% for Crypto. Weather, despite a small population ($n = 15$), runs a profitable share of 86.7% with median \$55K. Figures 20–21 visualise the contrast.




Figure 20: PnL distribution by primary category.




Figure 21: Per-category profitability and bot prevalence.

6.2 Why politics?

Three plausible mechanisms account for politics' dominance. **(i) Slow-resolving information:** political events feature daily news flow and high-quality polling that retail traders systematically under-weight. **(ii) Sticky priors:** partisan retail flow creates durable mispricings into which informed counter-flow can DCA over days or weeks. **(iii) Capital depth:** the largest political markets clear hundreds of millions in volume, supporting the size at which top wallets operate without market impact. Figure 22 shows the YES-bias distribution and its relationship to PnL. Median YES-bias is 46% (slightly NO-leaning), and there is no monotone relationship to profit—both deeply YES-biased and deeply NO-biased traders can profit, but the tail is thicker on the YES side, consistent with redemption-style strategies that hold YES tickets to a \$1 payoff.

Figure 22: YES bias and its relationship to profitability




Figure 22: YES position-share distribution and its link to PnL.

7 Behavioural Detail and Risk-Adjusted Performance

7.1 DCA intensity

Figure 23 plots average DCA entries per market against PnL, coloured by Smart Score. Median DCA is 15 entries; 28% of traders average more than 30 entries per market. Heavy-DCA wallets are over-represented in the right tail of profitability, but the relationship is non-monotone: extreme DCA (100 entries per market) also captures high-frequency bots that simply spam orders. The right interpretation is conditional—heavy DCA *combined with low sell rate and a politics specialization* is the highest-margin combination.




Figure 23: DCA intensity vs PnL.

7.2 Concentration vs diversification

The Polydata HHI scales from 0 (concentrated) to 100 (diversified). Figure 24 shows that profitable outcomes occur across the full HHI range; the very largest profits, however, cluster at

moderate diversification ($\text{HHI} \in [40, 80]$), which is consistent with betting a few related markets per political event rather than spreading across unrelated themes.




Figure 24: Position concentration (HHI) vs PnL.

7.3 Equity curves of the elite

Figure 25 overlays the daily cumulative PnL of the top-10 wallets by realised PnL. Several wallets show the staircase profile of discrete event resolutions; others compound smoothly. Drawdowns are shallow (most $< 20\%$ of peak; Figure 28 plots the cross-section). For winning traders in our sample, drawdown is *not* the binding constraint—most are still near their all-time-high cumulative PnL.




Figure 25: Equity curves for the top-10 PnL wallets.

7.4 Risk-adjusted performance

Figure 26 plots Sharpe vs PnL: most outsized profit producers sit in the modest-Sharpe (0.5–1.5) zone—i.e., they ran size and direction, not many uncorrelated bets. Figure 27 plots profit factor against trade win rate; the cross-shape highlights that high profit factors (>5) can come either from high win rates *or* from a small number of very large wins (resolution redemptions).




Figure 26: Sharpe vs Net PnL.




Figure 27: Profit factor vs trade win rate.




Figure 28: Drawdown from peak vs net PnL (winners only).

7.5 Badges and the Smart-Score–Bot-Score relationship

Figure 29 shows the prevalence of Polydata’s qualitative “badges” across traders. Profitable, Sniper and Diamond-Hands are the most common; Lottery Player and Bot are the most distinctive. Figure 30 reports the Spearman correlation among the key continuous metrics. The strongest pairwise correlation is between *event win rate* and *Smart Score* ($\rho \approx +0.7$); the strongest negative is between Bot Score and median trade gap ($\rho \approx -0.6$), as expected. Figure 31 crosses Bot Score with Smart Score and finds essentially no positive relationship: **being a bot does not make you smart, and being smart does not make you a bot**. The two scores measure orthogonal aspects of trader behaviour.




Figure 29: Badge prevalence across analyzed traders.




Figure 30: Spearman correlation of key trader metrics.




Figure 31: Bot Score vs Smart Score.

8 Discussion: What is the source of smart-money edge?

The empirical pattern across all six results sections converges on a single, consistent picture of how Polymarket smart money makes money.

It is not arbitrage. Cross-market arbitrage—e.g., between correlated weekly contracts on the same underlying—would manifest as high trade win rates ($\geq 80\%$ on small ticket sizes) and low position concentration. We see neither. Median trade win rate is 58%, and HHI is moderately concentrated. Top-volume bots that we suspect *do* arbitrage are precisely the wallets that earn the lowest median PnL.

It is not high-frequency market making. Pure HFT bots, defined operationally as wallets with ≥ 500 trades/day and median gap below 30 seconds, account for 11.7% of our sample but only a single-digit share of aggregate PnL. The infrastructure cost of maintaining tight quotes on Polymarket appears, on this evidence, not to be recouped by spread capture; the central limit order book is illiquid enough that adverse selection eats market makers alive.

It is also not insider trading. Polymarket’s outcomes are publicly resolvable; the most profitable markets (politics, sports) have no privileged information channel comparable to corporate fundamentals. The high event-level win rate is consistent with informed but not insider-grade priors.

It is resolution-edge holding. The signature combination—93% event win rate, 58% trade win rate, median 15 DCA entries per market, 29% of wallets with *zero* sell rate, and 98% profitability share among politics specialists—describes a strategy that should be familiar to value investors. The trader takes a probabilistic view (e.g., “Trump wins state X with $> 70\%$ probability”) and *accumulates* YES shares whenever the market price drifts below their fair value estimate. They lose individual trades when the market drifts further against them, but the inventory pays \$1 on resolution if their view is correct. Mathematically, this is closer to selling deep out-of-the-money puts than to scalping—a high-probability, asymmetric-payoff trade with a hard binary outcome at a known horizon.

The role of bots is execution, not edge. The dominant profitable cohorts (Heavy-DCA Holder; Likely-human fast traders) use automation as a tool—to spread orders across many small fills, to react quickly to news shocks, and to monitor the order book around the clock—but the underlying view is human-supplied. Pure bots, optimised for speed without a directional view, lose money on average. This suggests a clean labour-division frontier: *algorithms for microstructure, humans for macro view*.

Implications. The Polymarket smart-money playbook is extensible. Any binary outcome with (i) slow-resolving public information, (ii) retail directional flow that resists news, and (iii) deep enough order books to absorb staggered DCA entries should exhibit the same risk-reward profile. Politics is the canonical example. As the platform expands into financial-economic and geopolitical contracts, the same infrastructure—a well-founded prior plus patient inventory accumulation—ought to carry over.

9 Conclusion

We constructed a wallet-level cross-section of 273 analyzed Polymarket traders and decomposed their behaviour along 106 metrics. Profits are extremely concentrated (top decile = 52% of positive PnL); the median trader wins 93% of events but only 58% of trades; 67% of analyzed wallets are at least partially automated, yet pure bots earn the lowest median PnL of any class. Politics is the dominant profit pool (98% profitable), driven by slow-resolving news flow and durable retail directional bias. K-Means clustering identifies six strategy archetypes, of which “Heavy-DCA Holder” is the highest-yield. Read together, the data tell a coherent story: *smart money on Polymarket is not arbitrage and not insider trading; it is patient, opinionated capital that accumulates inventory at a discount to its own forecast and waits for binary resolution*. Bots provide execution rails; humans provide the forecasts.

Limitations. The Polydata sample is endogenously selected toward profitable, high-volume wallets; we therefore cannot speak to unconditional profitability across all Polymarket participants.

Bot classifications rely on Polydata’s proprietary detector, whose false-positive rate we cannot independently audit. Cluster labels are post-hoc descriptive; alternative K choices yield similar splits but different boundary cases.

Reproducibility. All scraping, parsing, analysis, and figure-generation code is provided in the `code/` directory. Raw HTML and parsed JSON for all 273 wallets are released; the flat panel `traders.csv` is a single-file replacement for the analysis here.

References

- Kenneth J. Arrow, Robert Forsythe, Michael Gorham, Robert Hahn, Robin Hanson, John O. Ledyard, Saul Levmore, Robert Litan, Paul Milgrom, Forrest D. Nelson, et al. The promise of prediction markets. *Science*, 320(5878):877–878, 2008.
- Pavel Atanasov, Phillip Rescober, Eric Stone, Samuel A. Swift, Emile Servan-Schreiber, Philip Tetlock, Lyle Ungar, and Barbara Mellers. Distilling the wisdom of crowds: prediction markets vs. prediction polls. *Management Science*, 63(3):691–706, 2017.
- R. M. Griffith. Odds adjustments by american horse-race bettors. *American Journal of Psychology*, 62(2):290–294, 1949.
- Charles F. Manski. Interpreting the predictions of prediction markets. *Economics Letters*, 91(3):425–429, 2006.
- Polydata. Polydata.pro: real-time analytics for polymarket. <https://polydata.pro>, 2026.
- Polymarket. Polymarket: prediction-market protocol on polygon. <https://polymarket.com>, 2025.
- Erik Snowberg and Justin Wolfers. Explaining the favorite-long shot bias: is it risk-love or misperceptions? *Journal of Political Economy*, 118(4):723–746, 2010.
- Richard H. Thaler and William T. Ziemba. Anomalies: parimutuel betting markets: racetracks and lotteries. *Journal of Economic Perspectives*, 2(2):161–174, 1988.
- Justin Wolfers and Eric Zitzewitz. Prediction markets. *Journal of Economic Perspectives*, 18(2):107–126, 2004.